

A 15x15 grid visualization of a learned policy. The grid shows a sequence of actions (red, blue, orange, purple) starting from the bottom-left corner and moving towards the top-right corner, illustrating a path through the state space.

